

Data Splits for ML



GROKKERS
AI FOR EVERYONE

Pre-reqs



Python



NumPy and PANDAS, SciPy,
Visualizations



Elementary stats and maths



Some preprocessing steps – may need
ML as well, for advanced topics

Background

Numeric Data: Preprocessing involves handling missing values, scaling to a similar range, and possibly normalizing the distribution.

Text Data: Common preprocessing steps include text cleaning (removing stop words, punctuation, etc.), tokenization, and vectorization (converting text into numerical form, such as TF-IDF or word embeddings).

Image Data: Techniques like resizing, normalization of pixel values, and data augmentation (creating variations of existing images) are often used.

Time Series Data: Dealing with temporal aspects, handling missing values over time, and creating lag features are important steps in preprocessing time series data.

Topics

About Data,
feature types,
tabular form

General
inspection of
data quality

Handling
duplicates in
data

Missing value
analysis
(2 parts)

Handling
Outliers

Cardinality
assessment

Encoding of
discrete data

Scaling and
Normalization

Handling
Skewed
Distributions

Data Imbalance
Handling

Data Splitting

Various types of datasets

1. Training Set:

- **Purpose:** Used to train the machine learning model.
- **Size:** Typically the largest subset (60-80% of the data).
- **Usage:** The model learns patterns, relationships, and features from this set.

2. Validation Set:

- **Purpose:** Used for hyperparameter tuning and model selection during training.
- **Size:** Smaller than the training set (usually 10-20% of the data).
- **Usage:** Helps prevent overfitting by fine-tuning model parameters without contaminating the test set.

3. Testing Set (or Test Set):

- **Purpose:** Reserved for evaluating the model's performance on unseen data.
- **Size:** Independent subset not used during training or validation (10-20% of the data).
- **Usage:** Provides an unbiased assessment of how well the model generalizes to new, unseen instances.

Model evaluation methods

Random test/input method
(predict)

Leave-One-out-Cross-
Validation method

- N , the number of data points in the set.
- N separate times, the estimator is trained on all the data except for one point and a prediction is made for that point.
- the average error is computed and used to evaluate the model.

Train-test split, **holdout
method**

- data set is separated into two sets, called the training set and the testing set.
- The function approximator fits a function using the training set only.
- Then the function approximator is asked to predict the output values for the data in the testing set

K-Fold cross-validation

- The data set is divided into k subsets, and the holdout method is repeated k times.
- Each time, one of the k subsets is used as the test set and the other $k-1$ subsets are used as training set.
- Then the average error across all k trials is computed.

EVALUATION – 1 (RANDOM INPUT TEST SAMPLE)

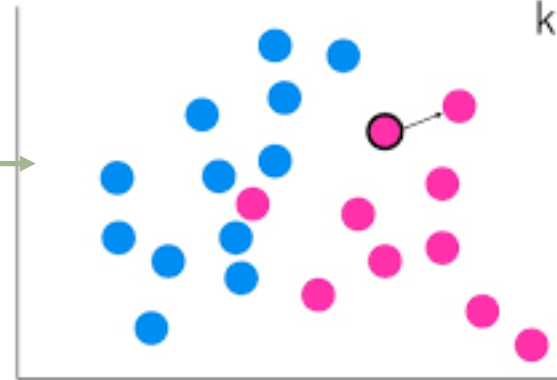
Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa
4.6	3.4	1.4	0.3	setosa
5.	3.4	1.5	0.2	setosa
4.4	2.9	1.4	0.2	setosa
4.9	3.1	1.5	0.1	setosa

Training dataset

KNN (k-nearest neighbors)

"what do similar cases look like?"

k=1



Test sample input



```
inputarray.append(input('first attribute : '))
inputarray.append(input('second attribute : '))
inputarray.append(input('third attribute : '))
inputarray.append(input('fourth attribute : '))
```

- Entire training set considered
- Test sample entered externally



	sepal_length	sepal_width	petal_length	petal_width	Class_name	distance
106	4.9	2.5	4.5	1.7	Iris-virginica	3.209361
114	5.8	2.8	5.1	2.4	Iris-virginica	3.416138
121	5.6	2.8	4.9	2.0	Iris-virginica	3.439477
59	5.2	2.7	3.9	1.4	Iris-versicolor	3.478505

EVALUATION – 2 (LEAVE-ONE-OUT CROSS-VALIDATION - LOOCV)

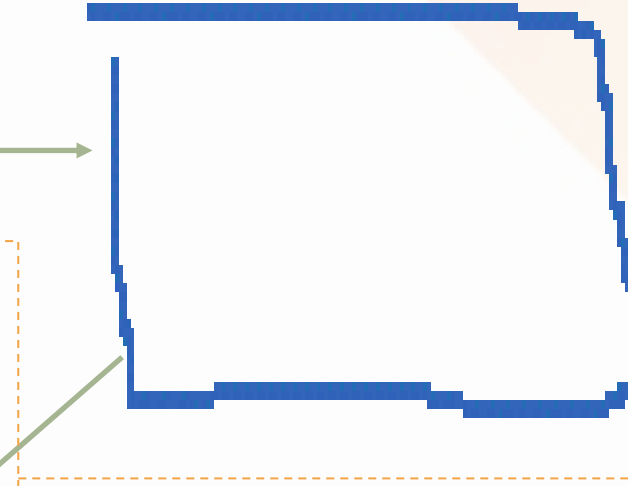
Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa
4.6	3.4	1.4	0.3	setosa
5.	3.4	1.5	0.2	setosa
4.4	2.9	1.4	0.2	setosa
4.9	3.1	1.5	0.1	setosa

Training dataset

Split into

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	NaN	setosa
2	4.7	NaN	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	NaN

Training



1 test sample

4	5.0	3.6	1.4	0.2	NaN
---	-----	-----	-----	-----	-----

- is a special case of cross-validation method in which each instance is used once as the test case and all other instances are used as the training set.
- also called as n-fold cross validation.
- utilizes the utmost training instances
- but due to its expensive nature it is usually applied to small datasets.



	sepal_length	sepal_width	petal_length	petal_width	Class_name	distance
106	4.9	2.5	4.5	1.7	Iris-virginica	3.209361
114	5.8	2.8	5.1	2.4	Iris-virginica	3.416138
121	5.6	2.8	4.9	2.0	Iris-virginica	3.439477
59	5.2	2.7	3.9	1.4	Iris-versicolor	3.478505

EVALUATION – 3 (HOLDOUT METHOD)

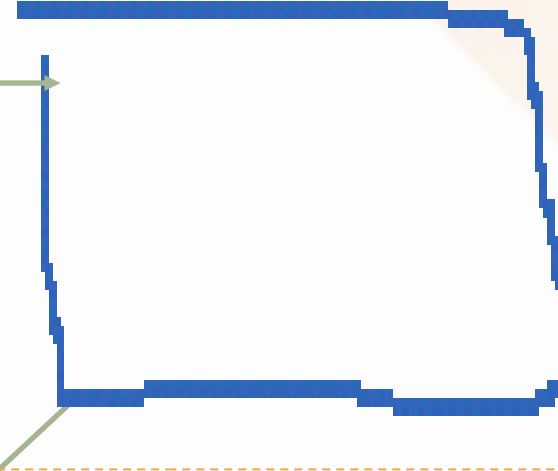
Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa
4.6	3.4	1.4	0.3	setosa
5.	3.4	1.5	0.2	setosa
4.4	2.9	1.4	0.2	setosa
4.9	3.1	1.5	0.1	setosa

Training dataset

Split into
Training
Test

sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2 setosa
1	4.9	3.0	1.4	NaN setosa
2	4.7	NaN	1.3	0.2 setosa
3	4.6	3.1	1.5	0.2 setosa

sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2 setosa
1	4.9	3.0	1.4	NaN setosa
2	4.7	NaN	1.3	0.2 setosa
3	4.6	3.1	1.5	0.2 setosa



Test sample

sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2 setosa
1	4.9	3.0	1.4	NaN setosa
2	4.7	NaN	1.3	0.2 setosa
3	4.6	3.1	1.5	0.2 setosa

- original data set is partitioned into two parts
- 50-50 or 70-30 or 60-40
- randomly divided into the training and test sets



	sepal_length	sepal_width	petal_length	petal_width	Class_name	distance
106	4.9	2.5	4.5	1.7	Iris-virginica	3.209361
114	5.8	2.8	5.1	2.4	Iris-virginica	3.416138
121	5.6	2.8	4.9	2.0	Iris-virginica	3.439477
59	5.2	2.7	3.9	1.4	Iris-versicolor	3.478505

Repeated for each test sample



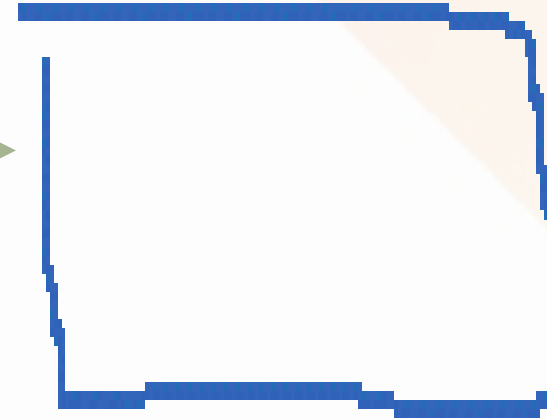
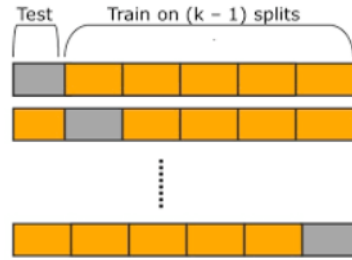
EVALUATION – 4 (K-FOLD CROSS-VALIDATION METHOD)

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa
4.6	3.4	1.4	0.3	setosa
5.	3.4	1.5	0.2	setosa
4.4	2.9	1.4	0.2	setosa
4.9	3.1	1.5	0.1	setosa

Training dataset

Split into

k-fold



Test sample

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	NaN	setosa
2	4.7	NaN	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa



Repeated for each fold

- Popular choices for K are 3, 5 and 10 as they're manageable computationally
- The cross-validation method is used with moderate datasets having instances around hundreds or more.



	sepal_length	sepal_width	petal_length	petal_width	Class_name	distance
106	4.9	2.5	4.5	1.7	Iris-virginica	3.209361
114	5.8	2.8	5.1	2.4	Iris-virginica	3.416138
121	5.6	2.8	4.9	2.0	Iris-virginica	3.439477
59	5.2	2.7	3.9	1.4	Iris-versicolor	3.478505



Repeated for each test sample

sklearn.model_selection: Model Selection

Splitter Classes	
model_selection.GroupKFold([n_splits])	K-fold iterator variant with non-overlapping groups.
model_selection.GroupShuffleSplit(...)	Shuffle-Group(s)-Out cross-validation iterator
model_selection.KFold([n_splits, shuffle, ...])	K-Folds cross-validator
model_selection.LeaveOneGroupOut()	Leave One Group Out cross-validator
model_selection.LeavePGroupsOut(n_groups)	Leave P Group(s) Out cross-validator
model_selection.LeaveOneOut()	Leave-One-Out cross-validator
model_selection.LeavePOut(p)	Leave-P-Out cross-validator
model_selection.PredefinedSplit(test_fold)	Predefined split cross-validator
model_selection.RepeatedKFold([n_splits, ...])	Repeated K-Fold cross validator.
model_selection.RepeatedStratifiedKFold(...)	Repeated Stratified K-Fold cross validator.
model_selection.ShuffleSplit([n_splits, ...])	Random permutation cross-validator
model_selection.StratifiedKFold([n_splits, ...])	Stratified K-Folds cross-validator
model_selection.StratifiedShuffleSplit(...)	Stratified ShuffleSplit cross-validator
model_selection.TimeSeriesSplit([n_splits, ...])	Time Series cross-validator
model_selection.train_test_split(*arrays, ...)	Split arrays or matrices into random train and test subsets

Demo using python/sklearn

(data splitters in
SCIKIT-LEARN)



Thanks !!

